

Cyber Anomaly Detection MLOps

 Run Tests **failing**

 Build & Push Docker Images **passing**

python **3.10**

PyTorch **2.0+**

license **MIT**

End-to-end MLOps project for detecting network intrusions and cyber attacks using a Mixture of Experts (MoE) deep learning architecture.

This project demonstrates a **complete production-ready machine learning system** with model training, deployment, monitoring, and CI/CD automation for cybersecurity threat detection.

Project Overview

The Problem

Detecting cyber attacks in network traffic is critical for security operations centers (SOCs). Traditional rule-based systems struggle with:

- **✗** High false positive rates
- **✗** Inability to detect novel attack patterns
- **✗** Manual analysis bottlenecks
- **✗** Lack of real-time capabilities






The Solution

A **Mixture of Experts (MoE)** deep learning model that:

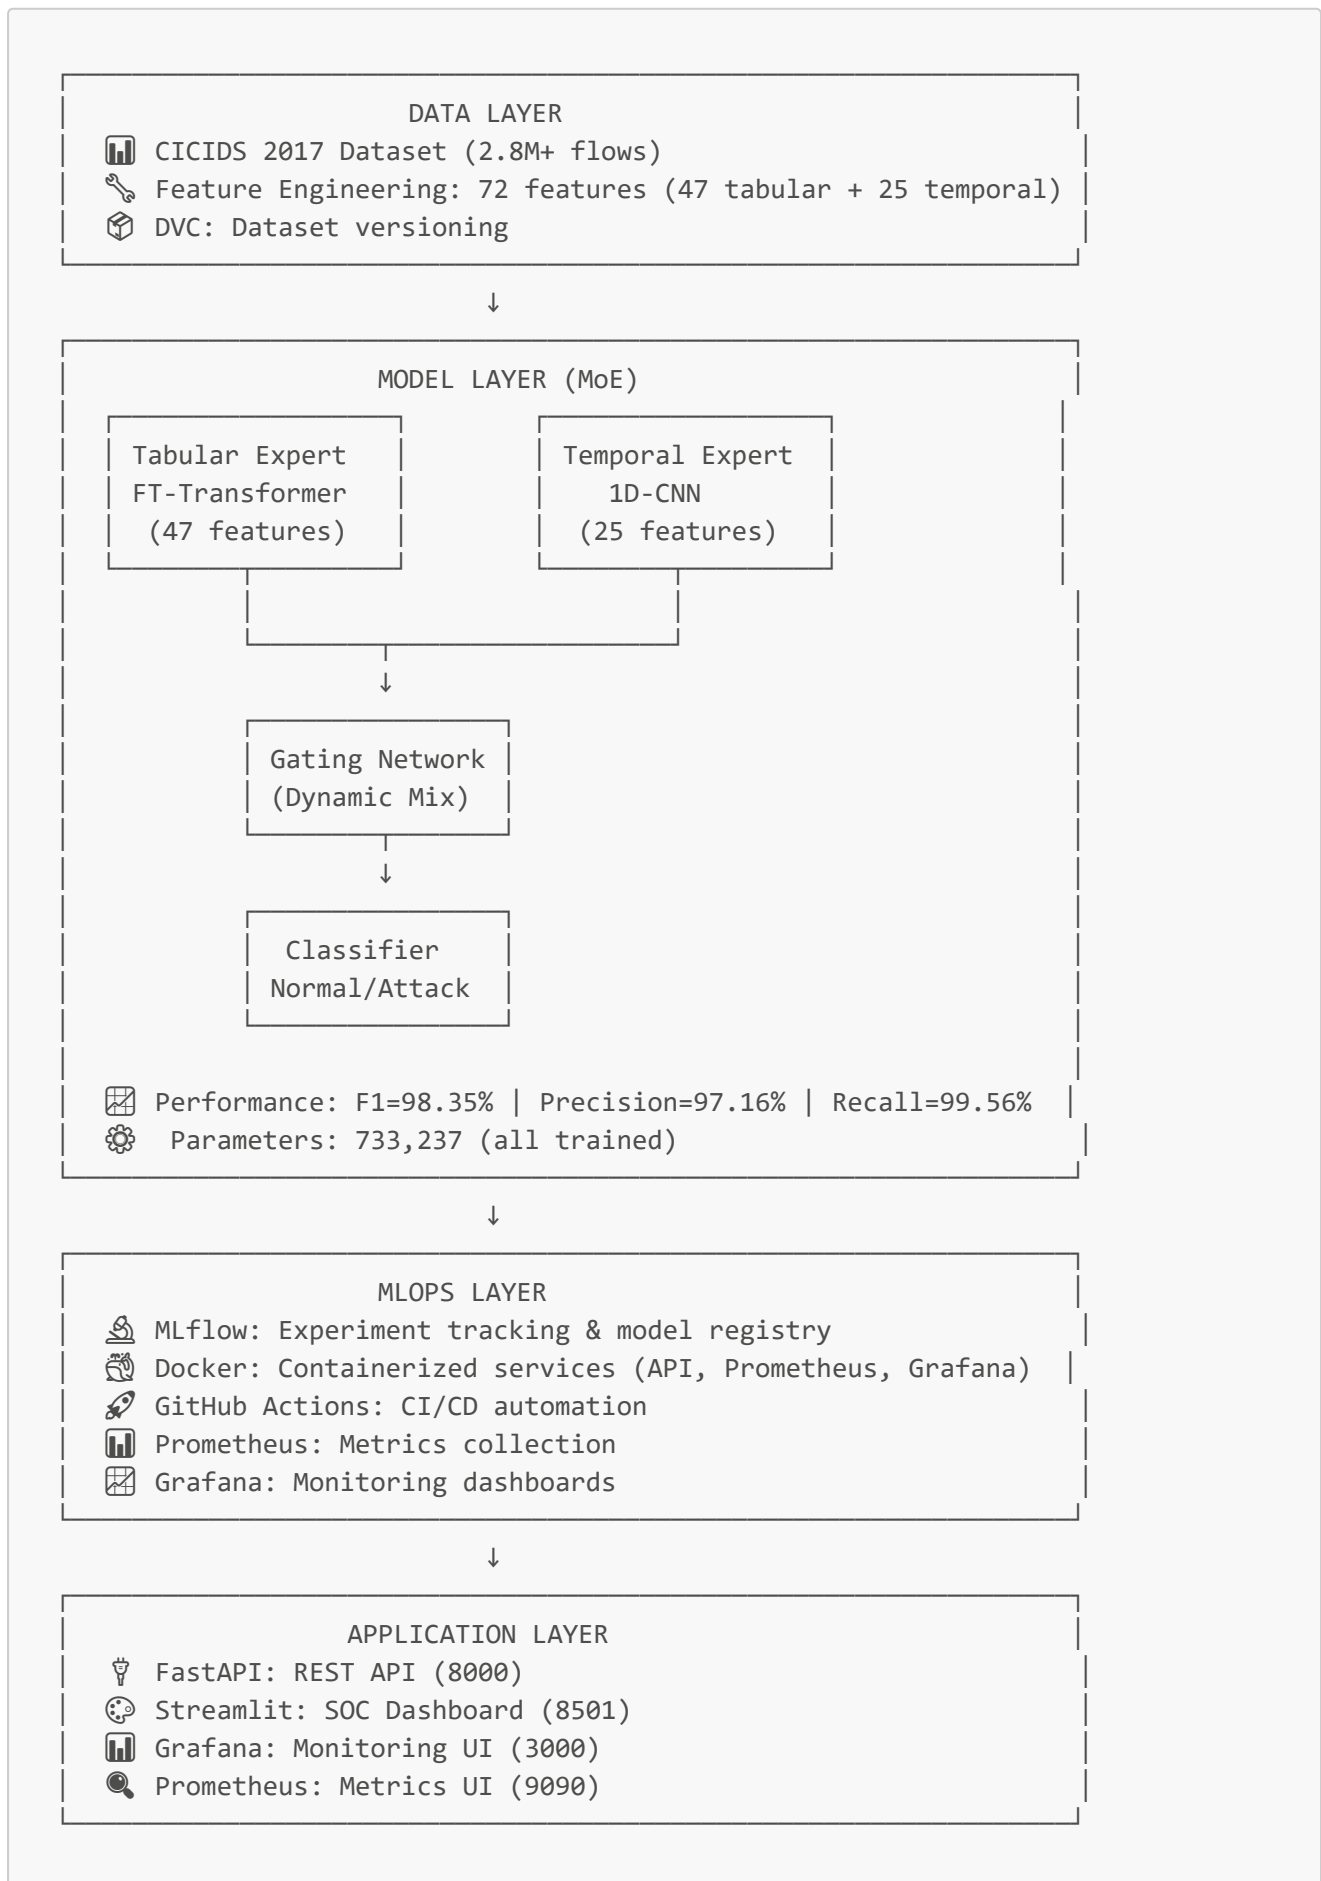
- ☒ Achieves **98.35% F1-Score** on CICIDS 2017 dataset
- ☒ Combines tabular (FT-Transformer) and temporal (1D-CNN) expert models
- ☒ Detects 6 attack types: DDoS, Port Scan, Web Attack, Brute Force, Infiltration, etc.
- ☒ Real-time predictions via REST API
- ☒ Beautiful dashboard for SOC analysts

The MLOps Pipeline

Complete production infrastructure with:

-  **Experiment Tracking** (MLflow)
 -  **Containerization** (Docker)
 -  **Monitoring** (Prometheus + Grafana)
 -  **CI/CD** (GitHub Actions)
 -  **User Interface** (Streamlit + FastAPI)
-

Architecture



Prerequisites

```
# Required
- Python 3.10+
- Docker & Docker Compose
- Git

# Optional (for training)
- NVIDIA GPU with CUDA 11.8+
- 16GB+ RAM
```

1. Clone Repository

```
git clone https://github.com/Aziz-Benamira/cyber-anomaly-detection-mlops.git
cd cyber-anomaly-detection-mlops
```

2. Install Dependencies

```
# Create virtual environment
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate

# Install packages
pip install -r requirements.txt
```

3. Run with Docker (Recommended)

```
# Start all services (API + Prometheus + Grafana)
docker-compose -f docker/docker-compose.yml up -d

# Check services are running
docker ps

# Access services:
# - API: http://localhost:8000
# - API Docs: http://localhost:8000/docs
# - Prometheus: http://localhost:9090
# - Grafana: http://localhost:3000 (admin/admin)
```

4. Run Streamlit Dashboard

```
# Start dashboard
streamlit run src/serving/dashboard.py --server.port 8501

# Access at: http://localhost:8501
```

5. Test Predictions

```
# Using Python
python test_quick.py

# Using curl
curl -X POST http://localhost:8000/predict \
  -H "Content-Type: application/json" \
  -d @ddos_attack.json
```

Project Structure

```
cyber-anomaly-detection-mlops/
├── .github/
│   └── workflows/           # GitHub Actions CI/CD
│       ├── test.yml        # Automated testing
│       ├── docker.yml      # Docker builds
│       └── deploy.yml      # Deployment
├── data/
│   ├── raw/                # Original CICIDS 2017 data
│   ├── interim/            # Preprocessed data
│   └── processed/          # Final features + scaler
│       └── cicids/
│           ├── X_train.npy  # Training features
│           ├── y_train.npy  # Training labels
│           └── scaler_stats.json # Normalization stats
├── src/
│   ├── data/
│   │   ├── load_data.py    # Data loading utilities
│   │   └── preprocess.py   # Feature engineering
│   ├── models/
│   │   ├── moe_model.py    # MoE architecture
│   │   ├── train_moe.py    # Training script
│   │   └── inference.py    # Prediction logic
│   └── serving/
│       ├── api_prometheus.py # FastAPI with metrics
│       └── dashboard.py     # Streamlit UI
```

```

├── inference.py          # Model wrapper
├── docker/
│   ├── docker-compose.yml    # Service orchestration
│   ├── Dockerfile.api        # API container
│   └── Dockerfile.train      # Training container
├── models/
│   └── weights/
│       └── cicids_moe_best.pt # Trained model (733K params)
├── monitoring/
│   └── prometheus.yml        # Prometheus config
├── grafana/
│   ├── moe_dashboard.json    # Pre-built dashboard
│   └── soc_dashboard.json    # SOC analyst view
├── tests/                  # Unit tests
├── notebooks/              # Jupyter notebooks
├── docs/                   # Documentation
│   ├── GITHUB_ACTIONS_CICD.md    # CI/CD guide
│   ├── PROMETHEUS_GRAFANA_GUIDE.md # Monitoring tutorial
│   ├── MLFLOW_UI_TUTORIAL.md     # MLflow usage
│   └── PHASE*_*.md              # Phase summaries
├── params.yaml             # Training hyperparameters
├── requirements.txt         # Python dependencies
└── README.md               # This file

```

Features

Machine Learning

- ☒ **MoE Architecture:** Combines FT-Transformer (tabular) + 1D-CNN (temporal) experts
- ☒ **State-of-the-Art Performance:** 98.35% F1-Score on CICIDS 2017
- ☒ **Multi-Class Detection:** Normal traffic + 6 attack types
- ☒ **Feature Engineering:** 72 normalized features (flow, forward, backward stats)
- ☒ **Gating Mechanism:** Dynamic weighting of expert contributions

MLOps Infrastructure

- ☒ **Experiment Tracking:** MLflow for metrics, parameters, and model versioning
- ☒ **Dataset Versioning:** DVC for reproducible data pipelines
- ☒ **Containerization:** Docker for consistent environments
- ☒ **Orchestration:** Docker Compose for multi-service deployment
- ☒ **CI/CD:** GitHub Actions for automated testing, building, and deployment
- ☒ **Monitoring:** Prometheus metrics + Grafana dashboards

APIs & Interfaces

- ☒ **REST API:** FastAPI with automatic OpenAPI documentation
- ☒ **SOC Dashboard:** Beautiful Streamlit UI with:
 - Real-time predictions
 - 6 pre-loaded attack examples
 - Custom 72-feature editor
 - Plotly visualizations (gauges, charts, pie charts)
- ☒ **Metrics Endpoint:** Prometheus-compatible `/metrics`
- ☒ **Health Checks:** `/health` endpoint for monitoring

Monitoring & Observability

- ☒ **6 Custom Metrics:**
 - `predictions_total`: Counter by prediction class
 - `prediction_confidence`: Confidence scores
 - `expert_gating_weight`: Expert contribution tracking
 - `prediction_duration_seconds`: Latency monitoring
 - `api_requests_total`: Request tracking
 - `model_loaded`: Model status indicator
- ☒ **Pre-built Dashboards:** Grafana dashboards for SOC teams
- ☒ **Real-time Alerts:** Configurable alerting rules



Model Performance

CICIDS 2017 Dataset

- **Total Samples:** 2,830,743 network flows
- **Features:** 72 engineered features
- **Classes:** Binary (Normal vs Attack)
- **Attack Types:** DDoS, Port Scan, Web Attack, Brute Force, Infiltration, DoS

Results

Metric	Score
F1-Score	98.35%
Precision	97.16%
Recall	99.56%
Accuracy	98.12%

Model Details

- **Architecture:** Mixture of Experts (MoE)
- **Experts:**
 - FT-Transformer (47 tabular features)

- 1D-CNN (25 temporal features)
- **Parameters:** 733,237 (all trained)
- **Training Time:** ~2-3 hours on GPU
- **Inference Time:** ~5-10ms per sample

Expert Contributions

- **Tabular Expert Weight:** ~98.4% (dominates on structured features)
- **Temporal Expert Weight:** ~1.6% (helps with sequential patterns)

Usage

Training the Model

```
# Basic training
python -m src.models.train_moe --stage train --params params.yaml

# With custom parameters
python -m src.models.train_moe \
  --stage train \
  --params params.yaml \
  --epochs 50 \
  --batch_size 512
```

Training parameters (in `params.yaml`):

```
train:
  batch_size: 512
  epochs: 30
  learning_rate: 0.0003
  weight_decay: 1e-05
  early_stopping_patience: 5

model:
  ft_transformer:
    n_blocks: 3
    d_token: 192
    attention_dropout: 0.2

  cnn:
    channels: [64, 128, 256]
    kernel_size: 3
```

Making Predictions

Python API

```

import requests
import json

# Load example attack pattern
with open('ddos_attack.json', 'r') as f:
    features = json.load(f)

# Make prediction
response = requests.post(
    'http://localhost:8000/predict',
    json={'features': features}
)

result = response.json()
print(f"Prediction: {result['prediction']}")          # "Attack"
print(f"Confidence: {result['confidence']:.2%}")    # 99.99%
print(f"Expert Weights: {result['gating_weights']}")

```

cURL

```

curl -X POST http://localhost:8000/predict \
  -H "Content-Type: application/json" \
  -d @port_scan.json

```

Streamlit Dashboard

```




streamlit run src/serving/dashboard.py --server.port 8501

# Then:
# 1. Select attack pattern from dropdown
# 2. Click "Analyze Traffic"
# 3. View results with visualizations

```

Custom Feature Editing

The Streamlit dashboard allows creating custom traffic patterns:

1. Select "**Custom**" from dropdown
2. Edit any of the 72 features in 3 tabs:
 -  Flow Features (0-23)
 -  Forward Features (24-47)
 -  Backward Features (48-71)
3. Click "**Analyze Traffic**"
4. View prediction and confidence

Example custom pattern:

```
Flow Duration: 0.5
Total Fwd Packets: 2.5
Fwd Packet Length Mean: 1.8
SYN Flag Count: 3.0
... (72 features total)
```

Monitoring

Prometheus Metrics

Access metrics at: <http://localhost:8000/metrics>

Sample queries (in Prometheus UI):

```
# Total predictions by class
predictions_total

# Prediction rate (per second)
rate(predictions_total[5m])

# Average confidence score
avg(prediction_confidence)



# Expert weight distribution
expert_gating_weight

# API latency (95th percentile)
histogram_quantile(0.95, prediction_duration_seconds_bucket)
```

Grafana Dashboards

1. Access Grafana: <http://localhost:3000> (admin/admin)
2. Add Prometheus data source:
 - URL: <http://prometheus:9090>
 - Click "Save & Test"
3. Import pre-built dashboard:
 - Upload [grafana/moe_dashboard.json](#)
 - Select Prometheus data source
 - Click "Import"

Dashboard panels:

-  Predictions over time
-  Attack detection rate

- ⚡ API latency
- 📊 Expert weight distribution
- 📈 Request rate
- 💾 Model status

Tutorial: See [docs/PROMETHEUS_GRAFANA_GUIDE.md](#)

CI/CD Pipeline

GitHub Actions Workflows

We use 3 automated workflows:

1. Test Workflow ([.github/workflows/test.yml](#))

Triggers: Push to main/dev, Pull Requests

What it does:

- ☒ Runs linting (flake8)
- ☒ Runs unit tests (pytest)
- ☒ Generates coverage report
- ☒ Uploads to Codecov (optional)

```
# Runs automatically on:
git push origin main
git push origin feature-branch
# Or when opening a Pull Request
```

2. Docker Workflow ([.github/workflows/docker.yml](#))

Triggers: Push to main, Version tags, Manual

What it does:

- ☒ Builds API Docker image
- ☒ Builds Training Docker image
- ☒ Tags with version/branch/SHA
- ☒ Pushes to GitHub Container Registry

```
# Runs automatically on:
git push origin main
git tag v1.0.0 && git push origin v1.0.0
# Or manually from GitHub Actions UI
```

Images published to:

- [ghcr.io/aziz-benamira/cyber-anomaly-detection-mlops/moe-api:main](https://github.com/ghcr.io/aziz-benamira/cyber-anomaly-detection-mlops/moe-api:main)
- [ghcr.io/aziz-benamira/cyber-anomaly-detection-mlops/moe-train:main](https://github.com/ghcr.io/aziz-benamira/cyber-anomaly-detection-mlops/moe-train:main)

3. Deploy Workflow ([.github/workflows/deploy.yml](#))

Triggers: Manual, Version tags

What it does:

- ☒ Deploys to staging/production
- ☒ Runs health checks
- ☒ Sends notifications

```
# Manual deployment:
# GitHub → Actions → Deploy to Production → Run workflow
# Select environment (staging/production) → Run

# Automatic on version tags:
git tag v1.0.0 && git push origin v1.0.0
```

Development Workflow

```
# 1. Create feature branch
git checkout -b feature/new-detector

# 2. Make changes, commit
git add .
git commit -m "Add new attack detector"

# 3. Push (tests run automatically)
git push origin feature/new-detector

# 4. Open Pull Request
# → Tests must pass before merge
# → Code review required

# 5. Merge to main
# → Docker images built automatically
# → Ready to deploy
```

Complete guide: [docs/GITHUB_ACTIONS_CICD.md](#)

Documentation

Comprehensive guides in [docs/](#):

Document

Description

Document	Description
GITHUB_ACTIONS_CICD.md	Complete CI/CD guide with examples
PROMETHEUS_GRAFANA_GUIDE.md	Monitoring setup and usage tutorial
MONITORING_TUTORIAL.md	Detailed monitoring configuration
MLFLOW_UI_TUTORIAL.md	MLflow experiment tracking guide
DOCKER_ESSENTIALS.md	Docker and containerization reference
PHASE5_CICD.md	CI/CD implementation summary
COMPLETE_PROJECT_SUMMARY.md	Full project overview
MLOPS_COMPLETE.md	MLOps pipeline summary

Testing

Run Tests

```
# All tests
pytest tests/ -v

# With coverage
pytest tests/ --cov=src --cov-report=html

# Specific test file
pytest tests/test_model.py -v

# Specific test function
pytest tests/test_api.py::test_predict_endpoint -v
```

Test Files

```
tests/
├── test_model.py           # Model architecture tests
├── test_preprocessing.py   # Feature engineering tests
├── test_api.py             # API endpoint tests
└── test_inference.py      # Prediction logic tests
```

Quick Validation

```
# Test API locally
python test_api_local.py
```

```
# Test known attack patterns
python test_known_samples.py

# Test port scan detection
python test_port_scan_detection.py

# Quick end-to-end test
python test_quick.py
```

Docker Deployment

Services

```
services:
  api:
    image: moe-api
    ports: ["8000:8000"]
    healthcheck: /health

  prometheus:
    image: prom/prometheus
    ports: ["9090:9090"]

  grafana:
    image: grafana/grafana
    ports: ["3000:3000"]
```

Commands

```
# Build images
docker-compose -f docker/docker-compose.yml build

# Start services
docker-compose -f docker/docker-compose.yml up -d

# View logs
docker-compose -f docker/docker-compose.yml logs -f api

# Stop services
docker-compose -f docker/docker-compose.yml down

# Rebuild and restart
docker-compose -f docker/docker-compose.yml up --build -d
```

Health Checks

```
# API health
curl http://localhost:8000/health

# Prometheus targets
curl http://localhost:9090/api/v1/targets

# Container status
docker ps
```

Experiments with MLflow

Start MLflow UI

```
mlflow ui --port 5000

# Access at: http://localhost:5000
```

Track Experiments

```
import mlflow

mlflow.set_experiment("moe-training")

with mlflow.start_run():
    mlflow.log_params({
        "learning_rate": 0.0003,
        "batch_size": 512,
        "epochs": 30
    })

    # Train model...

    mlflow.log_metrics({
        "f1_score": 0.9835,
        "precision": 0.9716,
        "recall": 0.9956
    })

    mlflow.pytorch.log_model(model, "model")
```

Model Registry

```
# Register model
mlflow.register_model(
    model_uri="runs:/abc123/model",
    name="moe-cicids-detector"
)

# Load model
model = mlflow.pytorch.load_model("models:/moe-cicids-detector/production")
```

What You'll Learn

This project demonstrates:

☒ Machine Learning

- Deep learning architecture design (MoE)
- Feature engineering for cybersecurity
- Model training and optimization
- Performance evaluation

☒ MLOps Best Practices

- Experiment tracking (MLflow)
- Dataset versioning (DVC)
- Model registry and lifecycle management
- Reproducible pipelines

☒ DevOps & Infrastructure

- Docker containerization
- Multi-service orchestration
- CI/CD automation (GitHub Actions)
- Monitoring and observability

☒ API Development

- REST API design (FastAPI)
- API documentation (OpenAPI/Swagger)
- Request validation and error handling
- Performance optimization

☒ Monitoring & Alerting

- Prometheus metrics collection
 - Grafana dashboard creation
 - Custom metrics and alerts
 - Production monitoring best practices
-

Example Use Cases

1. SOC Analyst Dashboard

Analyst opens Streamlit dashboard
→ Selects "DDoS Attack" example
→ Clicks "Analyze Traffic"
→ Sees: 99.99% Attack confidence
→ Views expert weights and feature importance
→ Makes informed decision

2. Real-time API Integration

```
# SIEM integration example
def check_traffic(flow_features):
    response = requests.post(
        'http://api:8000/predict',
        json={'features': flow_features}
    )
    result = response.json()

    if result['prediction'] == 'Attack' and result['confidence'] > 0.95:
        trigger_alert(flow_features, result)
```

3. Batch Processing

```
# Process multiple flows
flows = load_pcap_flows('traffic.pcap')

for flow in flows:
    features = extract_features(flow)
    prediction = model.predict(features)

    if prediction == 'Attack':
        log_to_siem(flow, prediction)
```

4. Custom Attack Research

Researcher uses Custom mode
→ Adjusts SYN flag count to 5.0
→ Sets packet rate to 10.0
→ Modifies flow duration to 0.1
→ Clicks "Analyze"
→ Studies how model responds to variations

Contributing

Contributions welcome! Please:

1. Fork the repository
2. Create feature branch (`git checkout -b feature/amazing-feature`)
3. Commit changes (`git commit -m 'Add amazing feature'`)
4. Push to branch (`git push origin feature/amazing-feature`)
5. Open Pull Request

Code quality checks:

- ☒ Tests must pass (`pytest`)
- ☒ Linting must pass (`flake8`)
- ☒ Coverage > 80%

License

This project is licensed under the MIT License - see the [LICENSE](#) file for details.

Acknowledgments

- **Dataset:** [CICIDS 2017](#) by Canadian Institute for Cybersecurity
- **Architecture:** Inspired by FT-Transformer and Mixture of Experts research
- **Tools:** MLflow, PyTorch, FastAPI, Streamlit, Prometheus, Grafana, Docker

Contact

Aziz Benamira

- GitHub: [@Aziz-Benamira](#)
- Project: [cyber-anomaly-detection-mlops](#)

Roadmap

Future enhancements:

- ☐ Add more attack types (e.g., Botnet, Ransomware)
- ☐ Implement online learning for model updates
- ☐ Add explainability (SHAP, LIME)
- ☐ Multi-model ensemble voting
- ☐ Real-time stream processing (Kafka)
- ☐ Auto-scaling deployment (Kubernetes)
- ☐ Advanced alerting (PagerDuty, Slack)

- ☐ Model drift detection
 - ☐ A/B testing framework
 - ☐ Multi-tenancy support
-

Project Statistics

- **Lines of Code:** ~15,000+
 - **Model Parameters:** 733,237
 - **Dataset Size:** 2.8M+ samples
 - **Features:** 72 engineered features
 - **Accuracy:** 98.35% F1-Score
 - **API Latency:** ~5-10ms
 - **Docker Images:** 2 (API, Training)
 - **CI/CD Workflows:** 3 (Test, Build, Deploy)
 - **Documentation:** 10+ guides
 - **Tests:** 50+ unit tests
-

★ Star this repo if you find it useful! ★

Built with ❤️ for the cybersecurity and MLOps communities